


# Segmentation by U-Net

**First Lecture: U-Net**


By Hongrun Zhang

# Semantic Segmentation


Only care about which category a pixel belongs to




Cell



X-Ray



Microscopic images



CT


# Instance Segmentation

Two levels:

1. A pixel belongs to which category?
2. A pixel belongs to which instance?




Semantic Segmentation

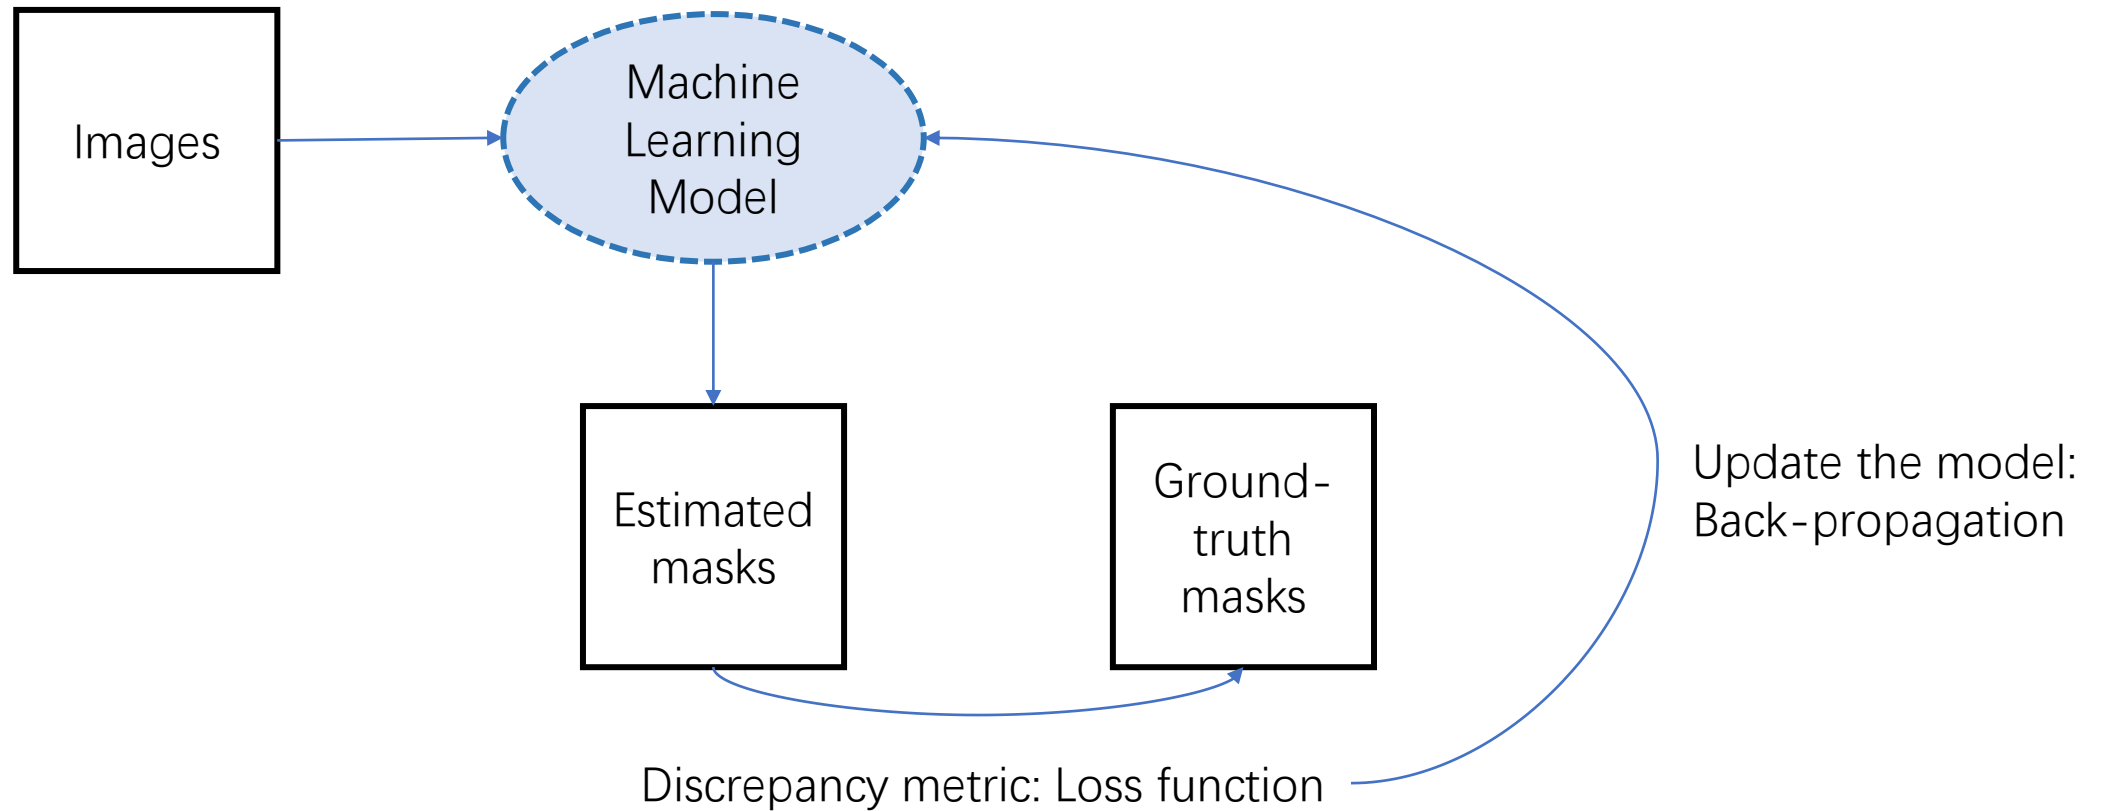


Instance Segmentation

# Machine learning-based semantic segmentation



# Machine learning-based semantic segmentation: Training process




# Loss functions

## Cross-entropy

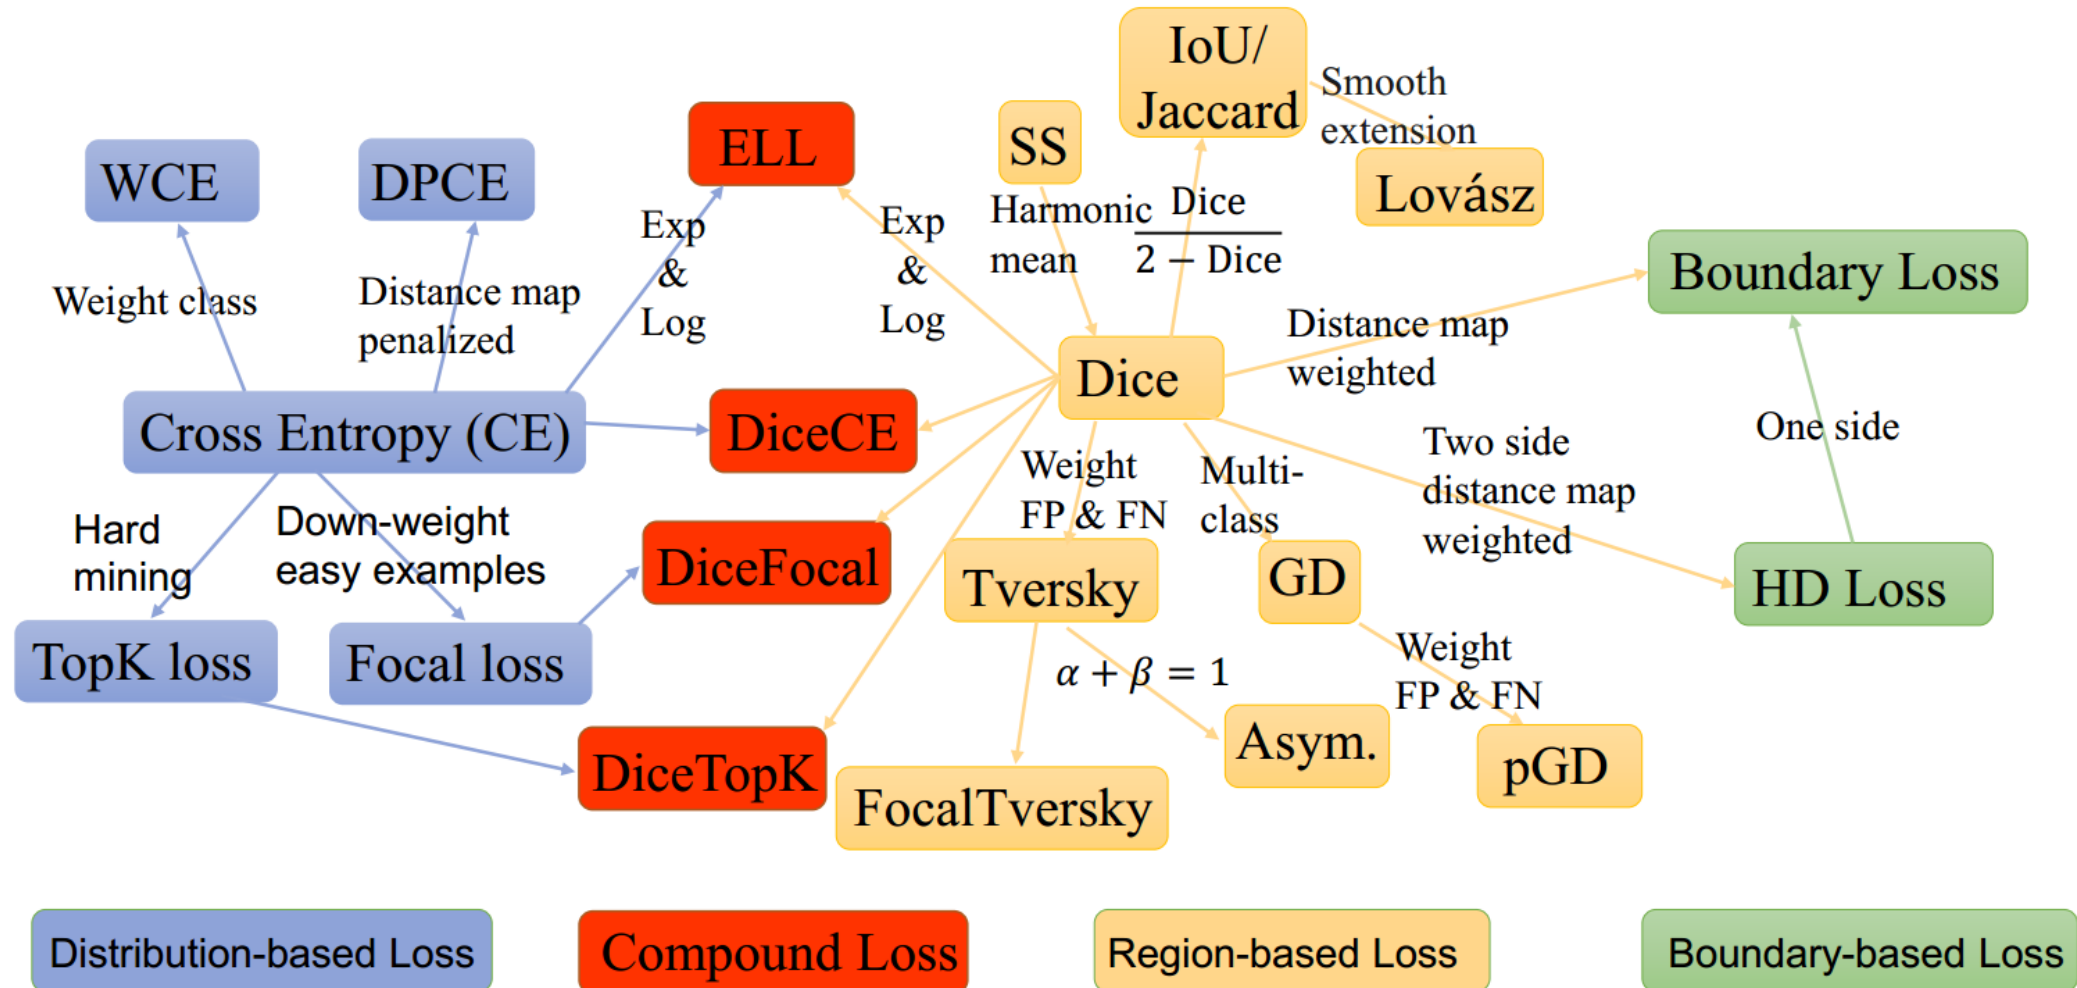
$$L(y, \hat{y}) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

## Dice coefficient

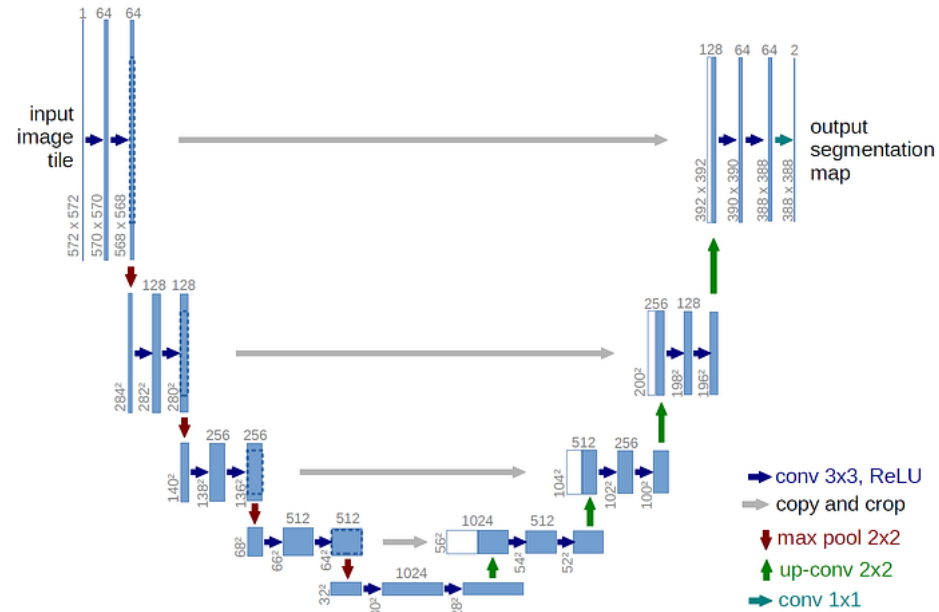
$$D = \frac{2 \sum_i^N p_i g_i}{\sum_i^N p_i^2 + \sum_i^N g_i^2}$$



# Loss functions




# U-Net



Basic Components


# Convolution Auto-encoder







# Basic convolution Module: 2D convolution

$$y[m, n] = x[m, n] * h[m, n] = \sum_{j=-\infty}^{\infty} \sum_{i=-\infty}^{\infty} x[i, j] \cdot h[m - i, n - j]$$



$$3 \times 1 + 2 \times 0 + 5 \times 1 + 2 \times 0 + 3 \times 3 + 4 \times 0 + 7 \times 1 + 1 \times 0 + 1 \times 1 = 7$$

# Basic convolution Module: 2D Pooling




# Basic convolution Module: Activation function

To introduce non-linearity for activation values


1. Sigmoid

$$f(x) = \frac{1}{1 + e^{-x}}$$




3. ReLU

$$f(x) = \max(0, x)$$




2. TanH

$$f(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$



4. Leaky ReLU

$$f(x) = \max(0.1x, x)$$

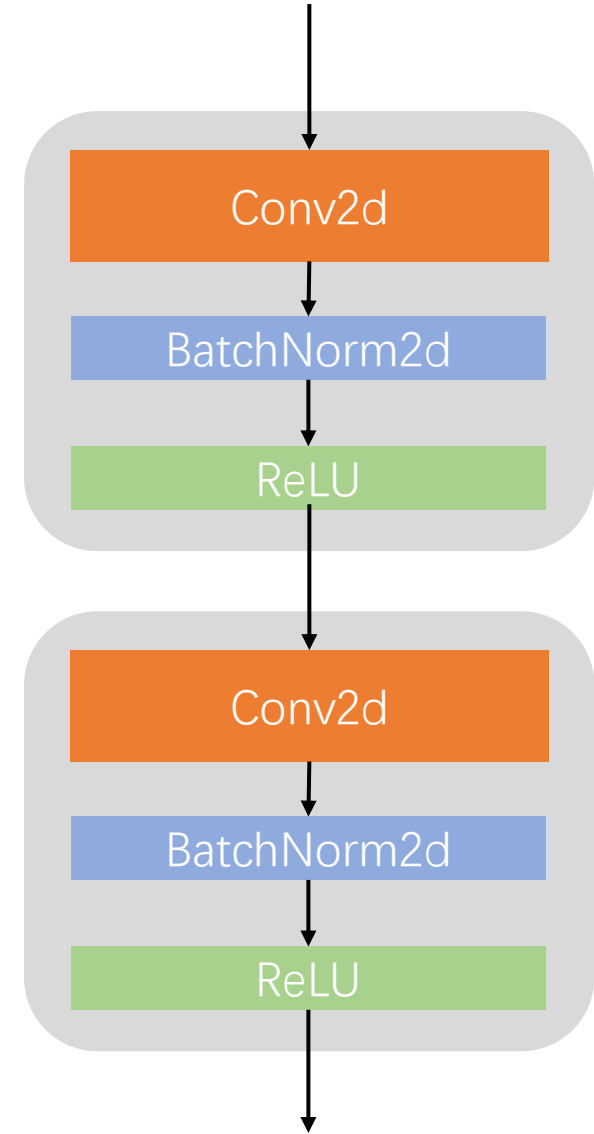


# Double-conv block

```
class DoubleConv(nn.Module):
    """(convolution => [BN] => ReLU) * 2"""

    def __init__(self, in_channels, out_channels, mid_channels=None):
        super().__init__()
        if not mid_channels:
            mid_channels = out_channels
        self.double_conv = nn.Sequential(
            nn.Conv2d(in_channels, mid_channels, kernel_size=3, padding=1, bias=False),
            nn.BatchNorm2d(mid_channels),
            nn.ReLU(inplace=True),
            nn.Conv2d(mid_channels, out_channels, kernel_size=3, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True)
        )

    def forward(self, x):
        return self.double_conv(x)
```




# Encoder: down-sampling module

```
class Down(nn.Module):
    """Downscaling with maxpool then double conv"""

    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.maxpool_conv = nn.Sequential(
            nn.MaxPool2d(2),
            DoubleConv(in_channels, out_channels)
        )

    def forward(self, x):
        return self.maxpool_conv(x)
```



# Decoder: Up-sampling module


```
class Up(nn.Module):
    """Upscaling then double conv"""

    def __init__(self, in_channels, out_channels, bilinear=True):
        super().__init__()


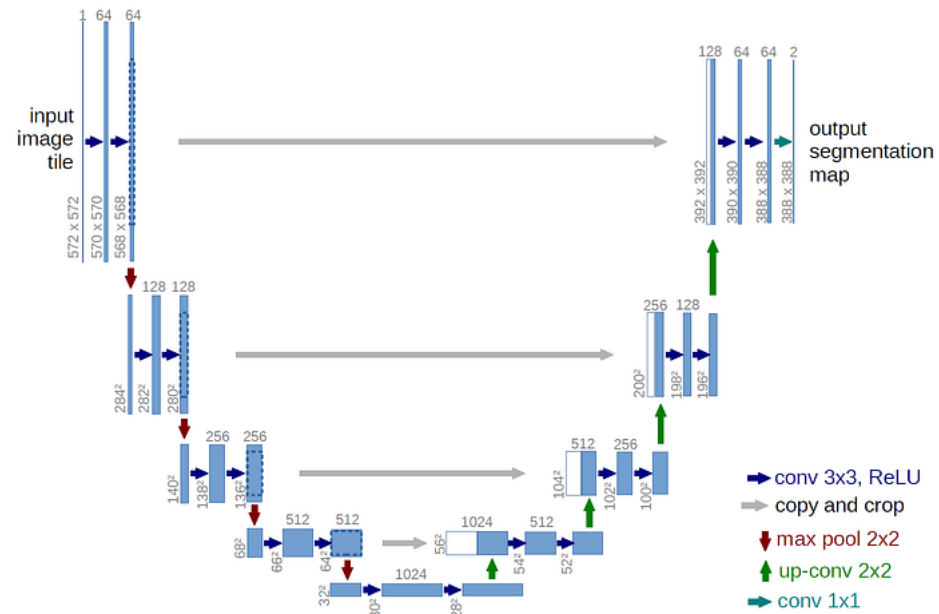
        # if bilinear, use the normal convolutions to reduce the number of channels
        if bilinear:
            self.up = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
            self.conv = DoubleConv(in_channels, out_channels, in_channels // 2)
        else:
            self.up = nn.ConvTranspose2d(in_channels, in_channels // 2, kernel_size=2, stride=2)
            self.conv = DoubleConv(in_channels, out_channels)

    def forward(self, x1, x2):
        x1 = self.up(x1)
        # input is CHW
        diffY = x2.size()[2] - x1.size()[2]
        diffX = x2.size()[3] - x1.size()[3]


        x1 = F.pad(x1, [diffX // 2, diffX - diffX // 2,
                        diffY // 2, diffY - diffY // 2])
        # if you have padding issues, see
        # https://github.com/HaiyongJiang/U-Net-Pytorch-Unstructured-Buggy/commit/0e854509c2cea854e247a9c615f175f76fbb2e3a
        # https://github.com/xiaopeng-liao/Pytorch-UNet/commit/8ebac70e633bac59fc22bb5195e513d5832fb3bd
        x = torch.cat([x2, x1], dim=1)
        return self.conv(x)
```



# Whole U-Net structure



# U-Net Variant



## U-Net Variant

Res-UNet

Dense-UNet

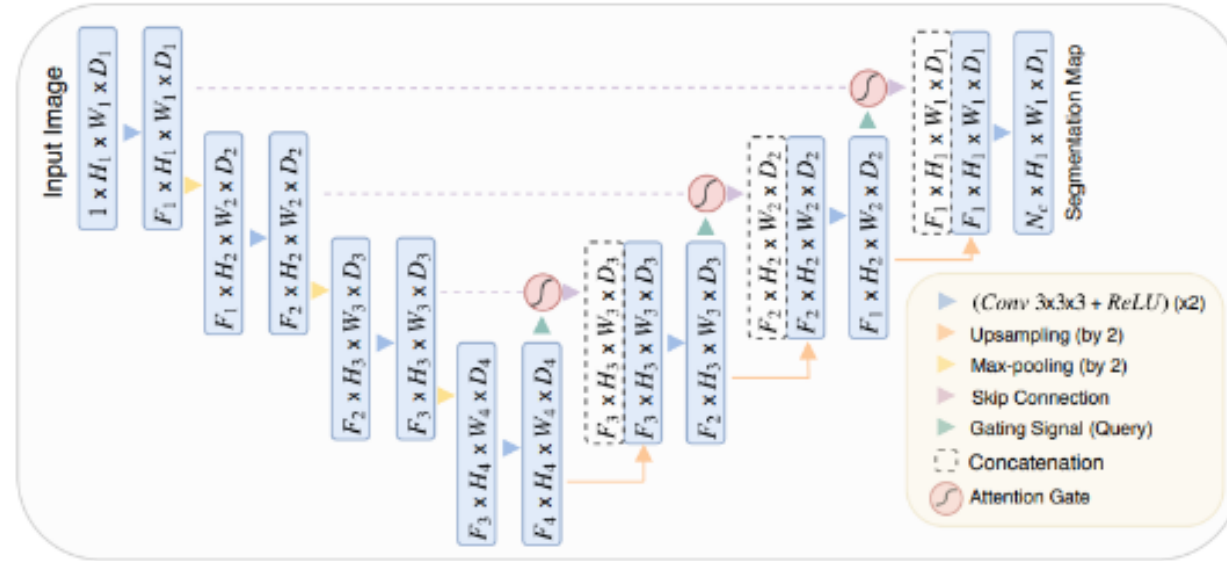
Attention-UNet

Unet++

Transformer-UNet




# Attention U-Net



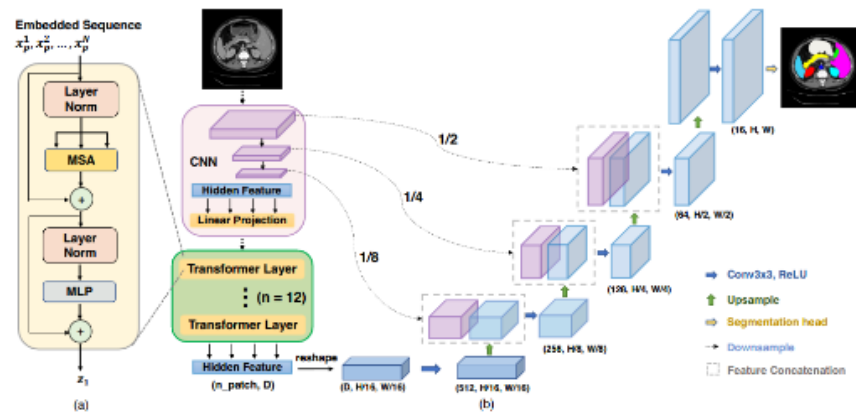
Ozan Oktay et al. 'attention u-net: Learning where to look for the pancreas attention'

# U-Net++



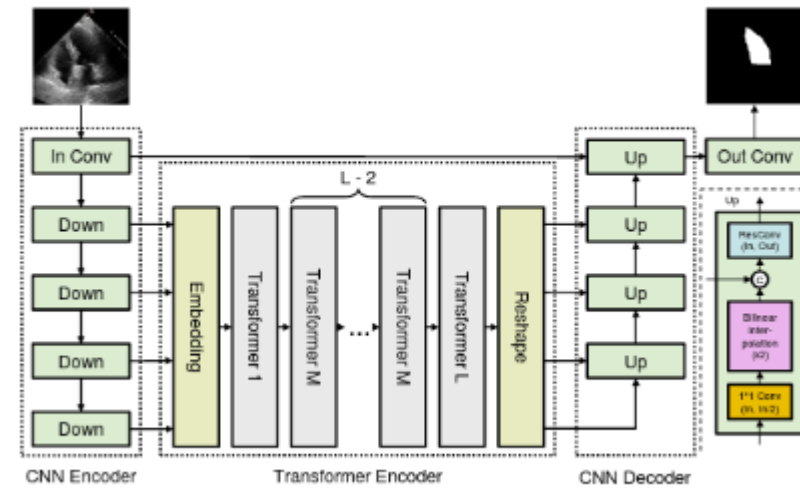
Zongwei Zhou. 'Unet++: A nested u-net architecture for medical image segmentation'

# Transformer-based U-Net



TransUNet

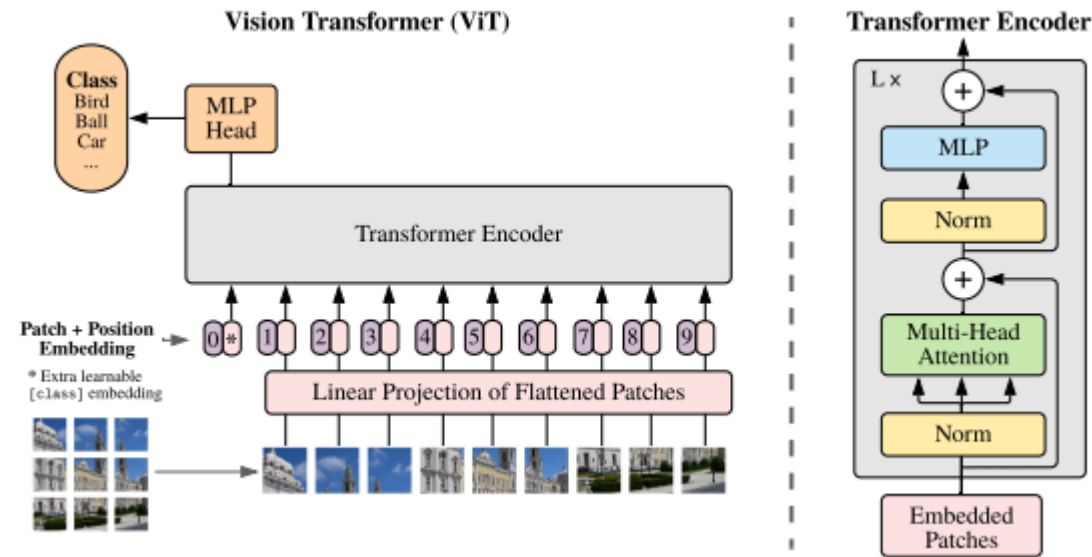
Jieneng Teng et al 'Transformers make strong encoders for medical image segmentation'



TransBridge

Kaizhong Deng et al 'Transbridge: A lightweight transformer for left ventricle segmentation in echocardiography'

# Vision transformer (ViT)



Dosovitskiy A, Beyer L, Kolesnikov A, et al. An image is worth 16x16 words: Transformers for image recognition at scale[J]. arXiv preprint arXiv:2010.11929, 2020.

# Task

- Run the code provided in github to come out with some segmentation results
- Try to replace the vanilla u-net with other architectures and redo the experiments. Compare the results.

[https://github.com/LeeJunHyun/Image\\_Segmentation](https://github.com/LeeJunHyun/Image_Segmentation)